



CS24: INTRODUCTION TO COMPUTING SYSTEMS

Spring 2014


Lecture 14

LAST TIME

- Examined several memory technologies:
 - SRAM – volatile memory cells built from transistors
 - Fast to use, larger memory cells (6+ transistors per cell)
 - DRAM – volatile memory cells built from capacitors
 - Slower to use, smaller memory cells, can make very large
 - Magnetic disk storage – nonvolatile memory
 - *Very* slow to use, compared to SRAM/DRAM/CPU!
 - Can make extremely large disks
- Also discovered two important principles:
 - Small memories tend to be faster than large ones, due to physical limitations
 - Can make larger memories denser (although addressing and accessing become more expensive)

LAST TIME (2)


- We want the performance of small memories, but we also want the convenience of large memories
- We will *never* get it with this design:



- The processor-memory gap:
 - Not only are processors much faster than large memories...
 - The gap itself has been growing by $> 50\%$ per year!
- Solution: use caching between CPU and memory

CACHING


- When CPU reads a value from main memory:
 - Read an entire block of data from main memory
 - As long as subsequent accesses are within this block, just use the cache!
 - If an access is outside the cached data, need to retrieve and cache a new block of data from memory



- Today:
 - What program behaviors maximize benefit of caching?
 - How are these caches designed? What are the trade-offs?

CACHING AND DATA ACCESS PATTERNS

- When CPU reads a value from main memory:
 - Read neighboring values from main memory as well



- Not all programs will benefit from this approach!
- Only programs that access data in specific patterns will benefit from the cache!
 - Most accesses end up being served from the cache
- Programs that usually access data in a widely varied manner will not benefit from cache at all.

CACHING AND LOCALITY

- Programs must exhibit good locality if they are going to utilize the cache effectively
- Spatial locality:
 - Program accesses data items that are close to other data items that have been recently accessed
- Temporal locality:
 - Program accesses the same data item multiple times
- Well-written programs will exhibit good locality
 - (“well-written” in terms of cache-friendliness)
 - ...and the computer hardware can run them faster!
- Poorly-written programs have poor locality
 - Program can’t take advantage of system caches

LOCALITY: EXAMPLES

- Frequently can achieve good locality very easily: programs tend to access data in regular patterns
- Vector-add code from before:

```
int * vector_add(int *a, int *b, int length) {  
    int i;  
    int *result =  
        (int *) malloc(length * sizeof(int));  
    for (i = 0; i < length; i++)  
        result[i] = a[i] + b[i];  
    return result;  
}
```

- Elements of input and output arrays are accessed in sequential order – works well with caching

LOCALITY: EXAMPLES (2)

- Still *extremely* valuable to understand locality as a programmer!
 - Simple choices can have a profound impact on program performance

- Molecular dynamics example from lecture 1

```
#define N_ATOMS 10000
#define DIM 2
/* Array of data for each atom being simulated. */
double atoms[N_ATOMS][DIM][DIM];
```

- Version 1:

```
for (i = 0; i < DIM; i++)
    for (j = 0; j < DIM; j++)
        for (n = 0; n < N_ATOMS; n++)
            atoms[n][i][j] = ... ;
```

- This code has poor data locality, so it runs slower

LOCALITY: EXAMPLES (3)

- Memory layout of our atoms array:

```
double atoms[N_ATOMS][DIM][DIM];
```

[0][0][0]	[0][0][1]	[0][1][0]	[0][1][1]	[1][0][0]	[1][0][1]	[1][1][0]	[1][1][1]
[2][0][0]	[2][0][1]	[2][1][0]	[2][1][1]	[3][0][0]	[3][0][1]	[3][1][0]	[3][1][1]

...

- Version 1:

```
for (i = 0; i < DIM; i++)  
  for (j = 0; j < DIM; j++)  
    for (n = 0; n < N_ATOMS; n++)  
      atoms[n][i][j] = ... ;
```

- Code accesses every 4th array element, because of the way the loops are arranged
- Rearrange loops to access each element in sequence

LOCALITY: EXAMPLES (4)

- Memory layout of our atoms array:

```
double atoms[N_ATOMS][DIM][DIM];
```

[0][0][0]	[0][0][1]	[0][1][0]	[0][1][1]	[1][0][0]	[1][0][1]	[1][1][0]	[1][1][1]
[2][0][0]	[2][0][1]	[2][1][0]	[2][1][1]	[3][0][0]	[3][0][1]	[3][1][0]	[3][1][1]

...

- Version 2:

```
for (n = 0; n < N_ATOMS; n++)  
  for (i = 0; i < DIM; i++)  
    for (j = 0; j < DIM; j++)  
      atoms[n][i][j] = ... ;
```

- This code accesses each array element in sequence
- A simple change, and results in a significant performance improvement!

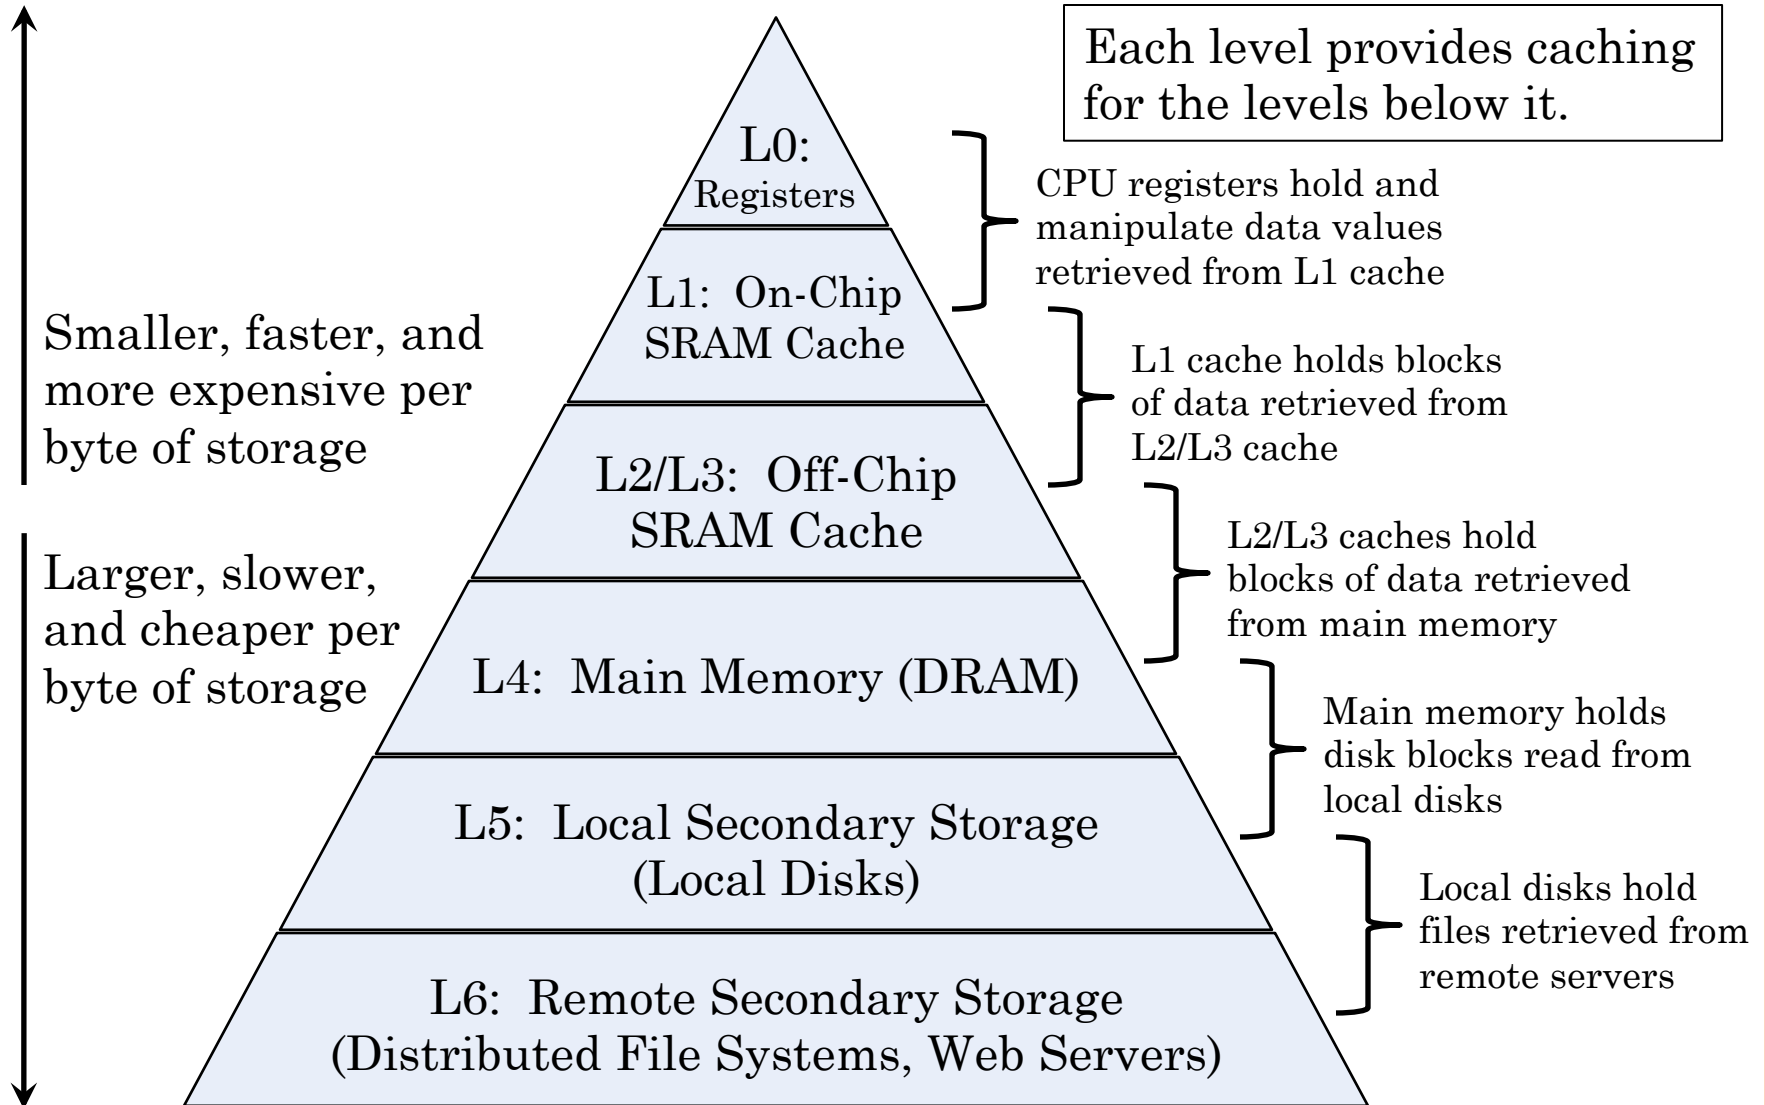
LOCALITY AND STRIDE

- When a program accesses array elements sequentially, called a stride-1 reference pattern
 - (Consider stride in terms of processor word-size)
 - Vector-sum program has a stride-1 reference pattern
 - Updated version of molecular dynamics program also has a stride-1 reference pattern
- When a program accesses every k^{th} element in sequence, called a stride- k reference pattern
 - Original version of molecular dynamics program has a stride-4 reference pattern
- Generally, as stride increases, program locality decreases
- With 1D arrays, pretty easy to achieve stride-1
- With multidimensional arrays, it can become much trickier to achieve stride-1

EXTENDING OUR CACHING IDEA

- Added an external SRAM cache between the CPU and main memory (DRAM)
- Accessing the SRAM cache is still slower than using registers...
 - Can access a register in 0 clocks (i.e. same clock that instruction is executed in)
 - SRAM cache can take e.g. 1-30 clocks, depending on size
- Also, disk access is horribly slow!
 - 20 million clocks or more!
 - Accessing DRAM is *much* faster than accessing the disk...
 - Could exploit data locality with disk accesses as well, by using DRAM as a cache for the disk
- *Why not apply our caching technique in other places?*

THE MEMORY HIERARCHY



THE MEMORY HIERARCHY (2)

- This is a very typical memory hierarchy used in modern computers, but by no means the only one!
- Many places where caching is employed within levels of the hierarchy
- Example: hard disks also employ caches
 - Disk buffer, often 2MB-64MB
 - Caches data prefetched from disk, pending writes to disk
 - Non-volatile RAM buffer (less common)
 - Disk writes are stored to this buffer until it fills up, then written to the disk itself
 - If power fails, this non-volatile memory retains its state
 - Finally, the magnetic disk storage itself
- Some systems use solid-state drives (SSDs) to cache data from traditional spinning hard-disks

CACHE MANAGEMENT

- Each level k provides caching for level $k + 1$
- Some caches are managed entirely by hardware
 - L1 (on-chip SRAM), L2/L3 (off-chip SRAM) caches
 - Performance is absolutely critical, so hardware is designed to manage these caches
- Other caches are managed entirely by software
 - L4 (DRAM main memory) is managed extensively by the operating system
 - e.g. the operating system caches disk blocks from L5 into L4 to improve disk IO performance
 - CPU registers (L0 cache) are manually assigned by compiler to minimize need to access memory (L1+)

CACHES AND MEMORY BLOCKS

- Level k caches data from level $k + 1$:
 - Memory at level $k + 1$ is partitioned into fixed-size blocks
 - Level k stores these blocks in its cache
 - Anytime data needs to be transferred between levels k and $k + 1$, this block size is used

- Actual block size depends on characteristics of levels k and $k + 1$

- Between L0 and L1, block size is 1 word
- Between L1 – L4, block is 8 to 16 words (64B)
- Between L4 and L5, block is up to several KB


L0	CPU Registers
L1	On-chip SRAM
L2/L3	Off-chip SRAM
L4	Main memory
L5	Local disk storage

- Lower levels have longer access times...
 - Also usually designed to read/write larger blocks of data from storage in one shot...

- Amortize read/write cost over a larger amount of data


CACHE LINES

- Hardware caches manage blocks of data from the next level...
 - Clearly need more details than just the data itself!
- Cache lines hold:
 - A flag indicating whether the line currently holds valid data
 - A tag that uniquely identifies the block
 - Taken from the address where the block is actually stored
 - The block of cached data itself
- Pictorially:



CACHE OPERATION


- CPU makes requests to main memory



- Main memory divided into blocks of B bytes ($B = 2^b$)
- Cache lines are slightly larger than B bytes
 - Line also includes the tag and a “valid” flag, as well as block
- Example: CPU requests a word in block 10
 - Cache returns value directly, since block 10 is cached
 - This is called a cache hit – the requested item was contained within the cache

CACHE MISSES


- Next, CPU requests a word in block 6



- Cache doesn't contain block 6
 - This is called a cache miss
 - Cache must load block 6 from main memory before providing requested value to the CPU
 - CPU incurs performance cost of accessing main memory
- What else must the cache do?
 - Figure out where to store block 6...
 - Need to *replace* (or *evict*) an existing block in the cache

REPLACEMENT POLICIES

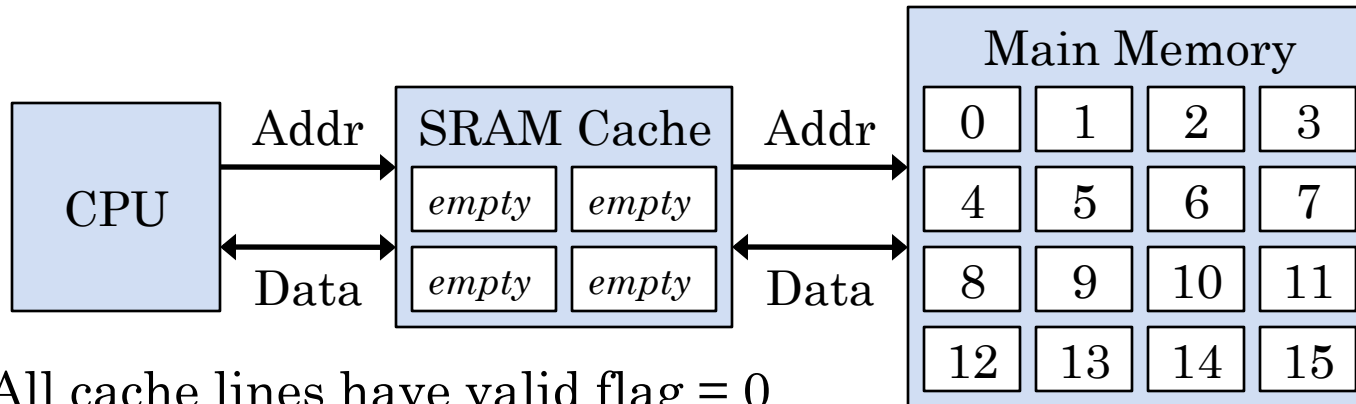
- CPU requested a value in block 6



- Caches have a replacement policy:
 - When loading a new block into a full cache, which existing block should be replaced?
- Example replacement policies:
 - Least Recently Used (LRU) policy: evict the block that was accessed furthest in the past
 - Random replacement policy: randomly pick a block to evict

COLD CACHES


- Will there be a cache miss in this situation?



- All cache lines have valid flag = 0
- Duh, of course!
- Caches have different kinds of cache misses
- This is called a *cold cache*
 - No accesses yet, so cache isn't populated with data!
 - Cache is “warmed up” by accessing memory and populating the cache lines
 - Misses during this phase are called *cold misses*
 - Also called “compulsory misses” since they are unavoidable

CACHE PLACEMENT POLICY

- CPU requests a word in block 5



- The biggest part of a cache's work?
 - Determining whether the requested block actually appears in the cache, and if so, where is it?!
- Caches implement a placement policy, specifying where new blocks are placed in the cache
- Most flexible placement policy is random
 - Blocks from level $k + 1$ can be stored anywhere in level k
- Also the most expensive placement policy!
 - Very costly to track down blocks in the cache
 - Hardware caches usually cannot implement this policy

CACHE PLACEMENT POLICY (2)

- Hardware caches use a much stricter placement policy
 - Blocks from level $k + 1$ can only be stored into a subset of locations in level k
 - Makes it much faster and easier to determine if a block is already in the cache
- Cache lines are grouped into cache sets
 - Every block from level $k + 1$ maps to *exactly one* set in the cache at level k
 - If cache set contains multiple cache lines, a block may be stored into any cache line in the set
- Two extremes for our cache organization:
 - S cache sets, each of which contains exactly one line
 - One cache set which contains all E cache lines

CACHE PLACEMENT POLICY (3)

- Cache lines are grouped into cache sets
 - Every block from level $k + 1$ maps to *exactly one* set in the cache at level k

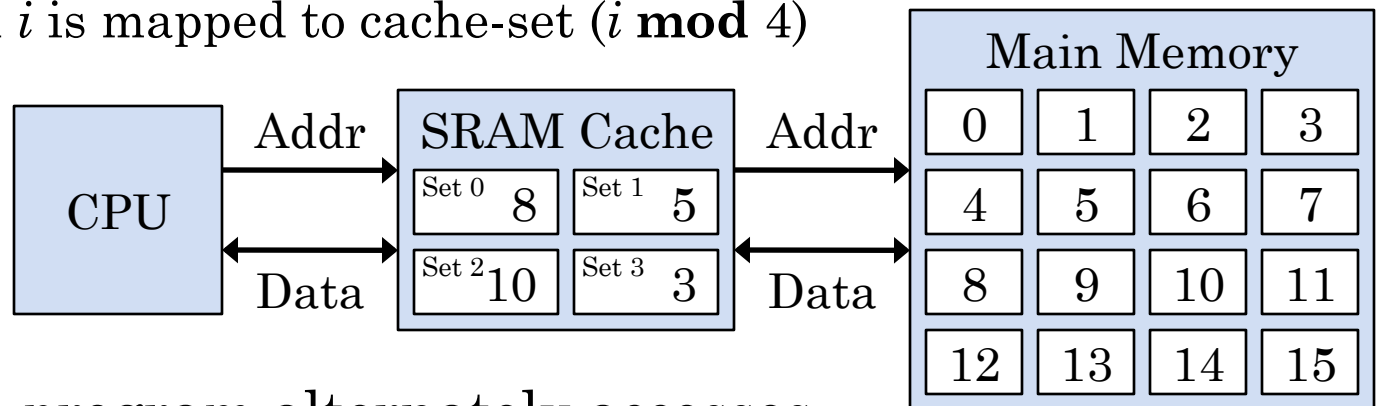
Cache line:



- Direct-mapped caches:
 - S cache sets, each of which contains exactly one line
 - Fast and easy to determine if a block is in the cache
 - Find the cache set associated with the block, and look at the one cache line in the set
- Fully associative caches:
 - One cache set which contains all E lines
 - Much more expensive to find if a block is in the cache
 - Need to look at the tags from *all* cache lines, to see if the block is in the cache

DIRECT-MAPPED CACHES AND MISSES


- Direct-mapped caches map each memory block to a single cache line
 - Each block only appears in one cache set, and each cache set only contains one cache line
- Can lead to new kinds of cache misses
- Example: our cache from before
 - Four cache sets, each of which contains one line
 - A block i is mapped to cache-set $(i \bmod 4)$



- What if a program alternately accesses data only in blocks 9 and 13?

DIRECT-MAPPED CACHES AND MISSES (2)

- Example: direct-mapped cache
 - A block i is mapped to cache-set ($i \bmod 4$)
 - Program alternately accesses data only in blocks 9 and 13

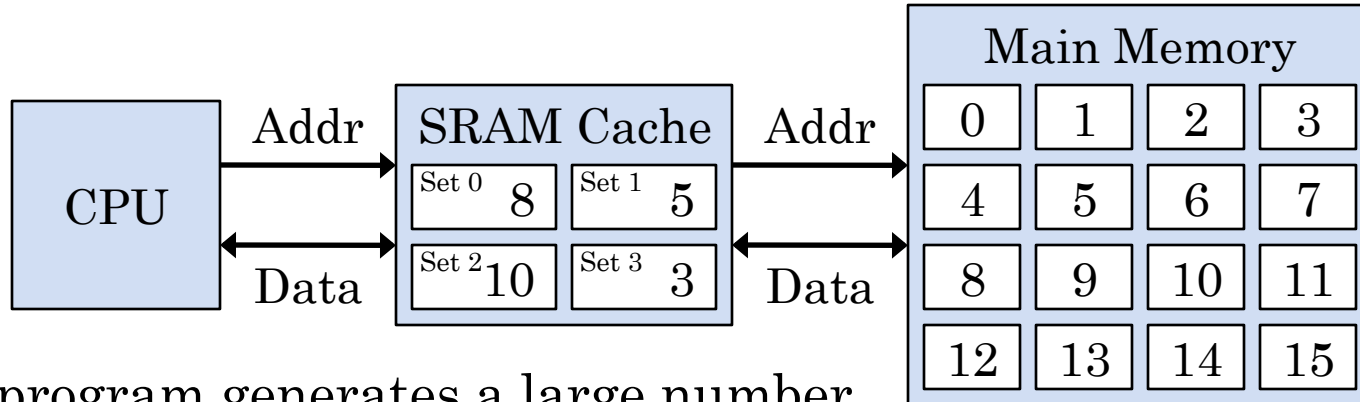


- Each time CPU makes a request, the cache doesn't contain the associated block!
- These are called conflict misses
 - Cache is large enough to hold the requested blocks...
 - ...but due to placement policy constraints, cache has to keep reloading blocks from main memory

DIRECT-MAPPED CACHES AND MISSES (3)

- Example: direct-mapped cache

- A block i is mapped to cache-set ($i \bmod 4$)
- Program alternately accesses data only in blocks 9 and 13



- If a program generates a large number of conflict misses like this, it is called thrashing
 - For our example, program thrashes between blocks 9 and 13
- ***A major issue!***
 - Program can have great locality, but still runs horribly slow!
 - If a program thrashes, the correct adjustments to make are often very subtle

CACHE ORGANIZATION


- Cache is designed to work against a main memory with M bytes
 - $m = \log_2(M)$ bits in addresses to main memory
- Caches have several important parameters
 - $B = 2^b$ bytes to store the block in each cache line
 - $S = 2^s$ cache sets
 - E cache lines per set
 - Both S and B are powers of 2
- The cache stores $B \times E \times S$ bytes of data from main memory
 - (Don't forget: cache lines also include a tag and a valid flag, which require additional space)

MAPPING CACHE BLOCKS


- Main memory with M bytes:
 - $m = \log_2(M)$ bits in addresses to main memory
- Cache parameters:
 - $B = 2^b$ bytes to store the block in each cache line
 - $S = 2^s$ cache sets
 - E cache lines per set
- Given a specific memory address, the cache must:
 - Map the address to a memory block
 - *(the cache works with blocks, not individual values)*
 - Figure out which cache set the block would live in
 - Figure out the tag that uniquely identifies the block
 - Figure out the offset of the address within the block
- Take m address bits, map them to these things

MAPPING CACHE BLOCKS (2)

- Relevant parameters for main memory and cache:
 - $m = \log_2(M)$ bits in addresses to main memory
 - $B = 2^b$ bytes to store the block in each cache line
 - $S = 2^s$ cache sets in the cache
- When CPU accesses a data value:



- Bottom-most b bits of the address specify offset of data value within the block




- Middle s bits specify the cache set where the block resides
- Remaining topmost bits constitute the block's tag
 - (Must be able to uniquely identify *all* blocks that can be cached)

MAPPING CACHE BLOCKS (3)

- Why use middle bits for set index?

- Why not topmost bits?



- If topmost bits identify the cache set, will cause long runs of addresses to map to same cache set
- Example: $M = 16$ bytes, $S = 4$ cache sets, $s = 2$
- If topmost bits select cache set, programs with good locality won't map accesses to different sets




- If middle bits select cache set, programs with good locality use all cache sets much more evenly



CACHE ACCESS OPERATIONS


- When cache receives a memory access, it must:
 - Figure out the cache set where the block goes
 - Figure out which cache line matches the block (*if any*)
 - Access the specific value within the cached block
- Given our mapping from memory addresses to cache details, these steps become pretty easy:



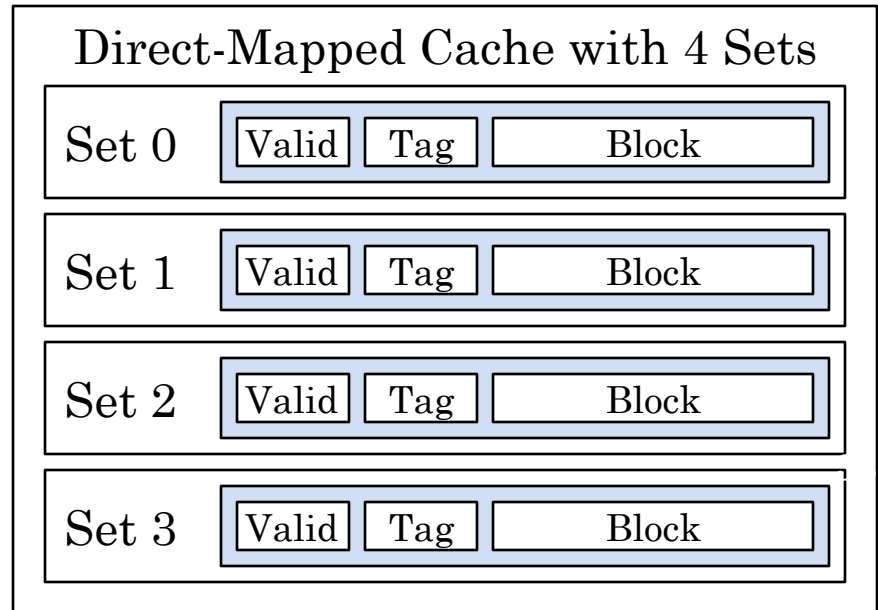
- Cache set's index is the middle s bits
- Use tag bits to identify the cache line within the set
- Bottom-most b bits are the access' offset within the block

DIRECT-MAPPED CACHES

- Direct-mapped caches have S sets, but each set contains only one cache line ($E = 1$)




- Example: direct-mapped cache with 4 sets
 - 2 bits in set index
- Very fast to map an address to a cache set
- Very fast to determine if a block is in the cache
 - If the tag doesn't match, block isn't in the cache!

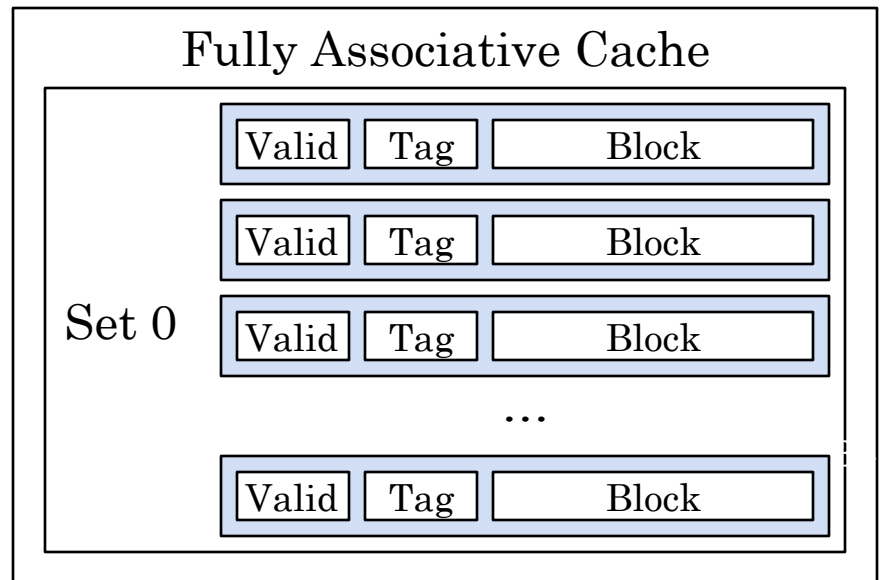


FULLY ASSOCIATIVE CACHES

- Fully associative caches have only one set, which contains all cache lines
 - $S = 2^s = 1 \Rightarrow s = 0$. No bits used for set index!



- Example: fully-associative cache
- Still very fast to map an address to a cache set
- More complicated to find if a block is in the cache
 - Need to examine all cache-line tags
 - Also, the tag is larger than in a direct-mapped cache




SET-ASSOCIATIVE CACHES

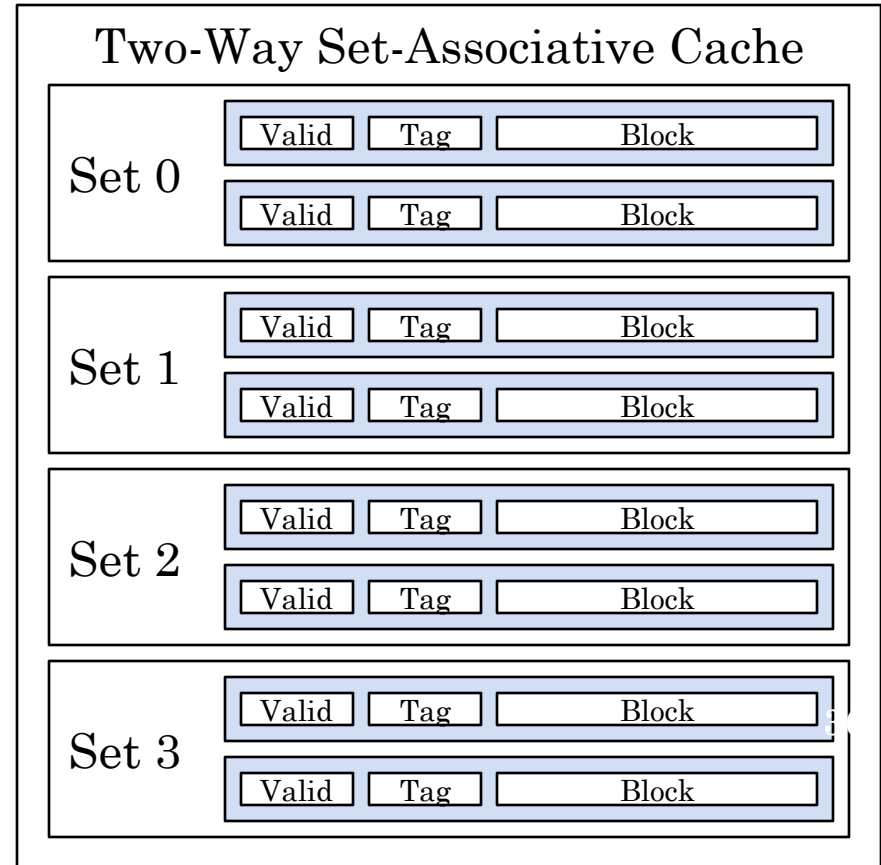
- Direct-mapped caches can easily cause a large number of conflict misses
- Fully associative caches must have complex logic to identify blocks by their tags
- Set-associative caches combine capabilities of both approaches into a single cache
 - Employ S cache sets, where $S > 1$
 - Each cache set contains E cache lines, where $E > 1$
- Achieves benefits of both techniques:
 - Greatly reduces potential for conflict misses
 - Limits the complexity of the logic that has to match block tags

SET-ASSOCIATIVE CACHES (2)

- Example: a two-way set-associative cache
 - E is number of cache lines in each set...
 - Cache is called “ E -way set-associative cache” when $S > 1$
 - $E = 2$ for this example
- For an m -bit address:

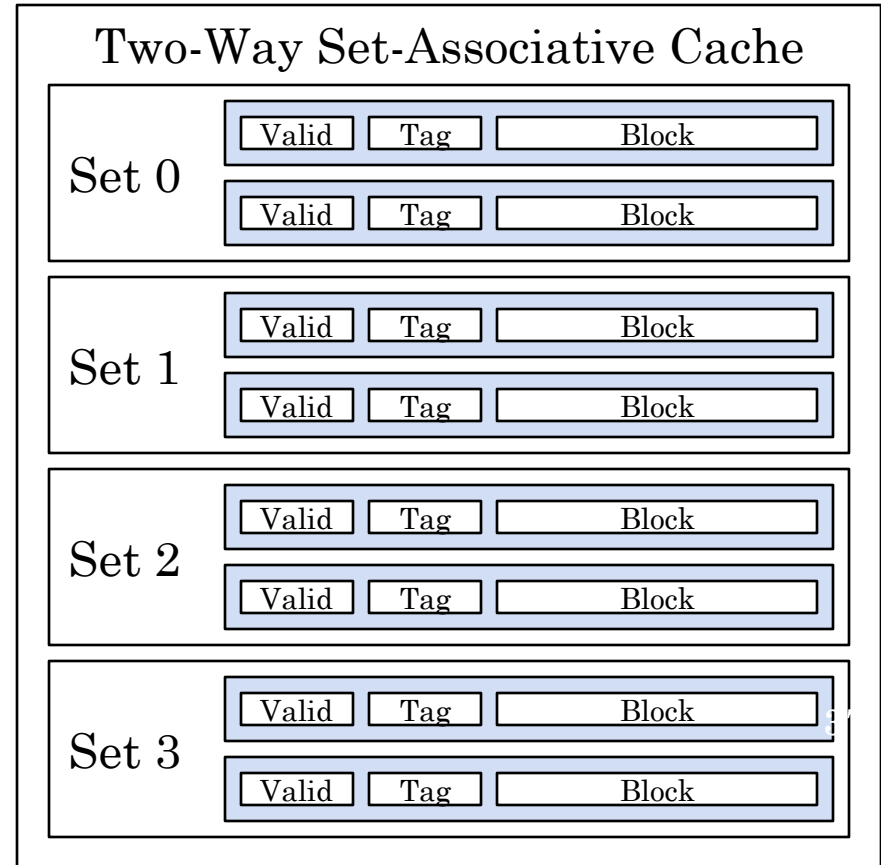


- Use set index to find cache set to examine
- Only have to check tags on small number of lines



LINE REPLACEMENT POLICIES

- When a cache miss occurs, must load a new block into the cache
 - Store in some cache line
- For direct-mapped caches, this is easy
 - Only one line per set
- For set-associative and fully associative caches, we get to choose a line!
- Replacement policy controls which cache line to evict when new block is loaded into cache

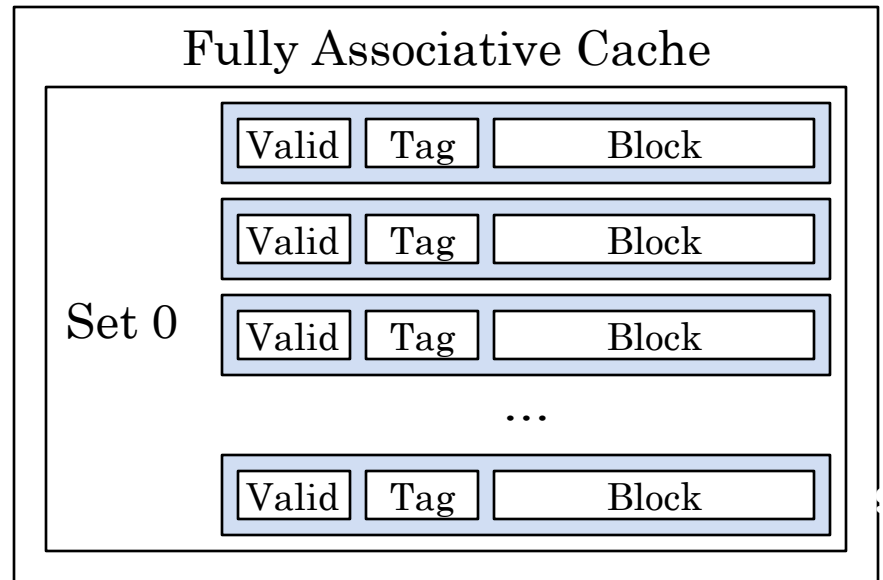


LINE REPLACEMENT POLICIES (2)

- Least Recently Used (LRU) policy
 - Evict cache line that was accessed furthest in past
- Least Frequently Used (LFU) policy
 - Evict cache line that was accessed the least frequently, over some time window
- These policies take extra time and hardware to implement
 - Not used as much in caches close to the CPU, where performance is critical
 - Used very often in caches further from the CPU, where cache misses are extremely costly
- For example, disk-block caches benefit *greatly* from more sophisticated replacement policies
 - *...when a cache miss costs 20 million clocks, spend a few thousand clocks to figure out what to keep in the cache...*

ASSOCIATIVE CACHES

- Where does the “associative” come from in set-associative caches and fully-associated caches?
- Each cache set has E cache lines in it...
 - Need to look up cache line using only the block's tag
 - The cache set is an *associative memory*
- Associative memory:
 - Not accessed with an address, like normal memories!
 - Associative memory stores (key, value) pairs
 - Key is the input to the associative memory
 - Memory returns value



ASSOCIATIVE CACHES (2)

- Associative caches must effectively implement associative memories for their cache sets
 - Keys are a concatenation of the tag, *plus* the valid flag
 - *No reason to look at the cache line if it isn't valid...*
 - Value is the block of data in the cache
- Set-associative caches:
 - Each cache set is an associative memory
 - Number of cache lines in each set is small, so logic is easier to implement
- Fully associative caches:
 - Need to examine *many* cache lines in parallel
 - *Much* more expensive...

